

Second Nature

A job-search agent that learns
from every decision

3-hour MVP



The repeated work in a job search

**“I already told you
why that role was wrong.”**

Illustrative job-seeker scenario

Another shortlist

Recheck seniority,
location
and the actual work.

Another application

Repeat facts already
confirmed elsewhere.

Another correction

Explain the same
preference
one more time.

A crowded market, a specific product bet

Product	Publicly described strength
Simplify	Application autofill and answer reuse
Teal	Tailored resumes and job tracking
Jobright	Profile-based matching and match scores

Our bet: auditable learning across repeated runs

Visible memory, workflow replay and user-controlled autonomy

The product loop

01

Verify

Resume facts and
standard answers

02

Judge

Rank roles and
explain the match

03

Learn

Keep or skip, then
confirm a preference

04

Reuse

Prepare a pack with
known answers

“This role fits your verified skills and avoids the seniority you rejected.”

Illustrative explanation. Final submission stays with the job seeker.

Real online jobs, fetched live

2,321


open technical roles

19 companies

Greenhouse, Ashby, Lever

Fetches Sep 11, 2026

US, Canada and remote



Five technologies, five distinct jobs

Layer	Technology	Concrete output
Structure	Cognee	Claims and preferences with provenance
Memory	HydraDB	Candidate graph and decision history
Insight	Hotdata	Rankable jobs and run-level analytics
Motion	RocketRide	Sequenced tool calls and application packs
Replay	Rote / Modiqo	Recorded tool paths for repeat work

Snyk checks dependencies and code before the demo.

The technical choices behind the loop

A fixed scorer, richer memory

45%

Verified coverage

30%

Retrieval relevance

25%

Preference fit

Rote: exact replay, partial replay or a new path

Cited claims in every pack. Autonomy requires user approval.

Evidence we can stand behind

85

offline tests passed

2,321 jobs
loaded and read back

Live compounding evidence
is the next demo gate.

Measure	Record on every run
Cost	Tokens or a labeled proxy
Human effort	Questions and approvals
Judgment	Recall on human-skipped jobs

Cost, effort and judgment together

A three-hour MVP scope

Time	Build checkpoint
0–30 min	Freeze input and record a baseline
30–90 min	Connect the five sponsor paths
90–145 min	Repeat runs with real human feedback
145–180 min	Scan, capture traces and rehearse

One candidate. One digest. One application pack. One measured loop.

An initial market in tech job search

10

design partners to recruit for the pilot

Developer communities and career coaches

Test a paid subscription during an active search

Success: less preparation time and shortlists people keep

Demo runbook: one visible learning cycle

Beat	What the audience sees
Baseline	A shortlist and one supported application pack
Feedback	A repeated skip reason becomes a proposed rule
Confirmation	The human approves it and the order changes
Replay	Known answers return, while a missing fact stays a question
Proof	The new run appends observed metrics

Real jobs. Assigned release days. Actual run measurements.

Technical boundaries and judging checks

Boundary	What we say and show
Data freshness	Live online collection; fetch timestamps retained
Graph execution	Cloud persistence; optional Bolt for server-side Cypher
Sponsor integration	Repeated live traces, not SDK imports or offline tests
Security	Snyk report, cited claims and human submission
Network	No board fetch on stage; cloud tools still need connectivity